

Static Symmetric Searchable Encryption over Attributes in SQL Databases

In this work we consider the following scenario. A user \mathbf{U} wants to store a collection of confidential documents $\mathbf{D} = \{D_0, \dots, D_m\}$ at untrusted server \mathbf{S} . To preserve data confidentiality, \mathbf{U} encrypts \mathbf{D} to obtain $\mathbf{C} = \{C_0, \dots, C_m\}$, which is outsourced to \mathbf{S} in such a way that 1) \mathbf{S} will learn as less as possible useful information about \mathbf{D} ; and that 2) \mathbf{S} can be given the ability to search through the collection and return appropriate (encrypted) documents to \mathbf{U} . We consider \mathbf{S} stores encrypted documents using relational data model and \mathbf{U} is able to use well defined SQL language to query \mathbf{S} in order to upload or retrieve encrypted documents.

A traditional symmetric searchable encryption mechanism allows searching the keywords directly in \mathbf{C} , without compromising data confidentiality nor query privacy. Searchable keywords are selected by \mathbf{U} and are part of a dictionary $\Delta = \{w_0, \dots, w_d\}$ of d unique words ordered lexicographically. Considering relational data model, we can say that single document $D_j = \{w_l^{0,j}, \dots, w_l^{n,j}\}$, is a j -th record in relation, while a single keyword $w_l^{i,j} \in \Delta$, $l \leftarrow \{1, d\}$, $i \in \{0, n\}$, $j \in \{0, m\}$, labeled by i -th attribute and j -th record is an attribute value (see Figure 1).

	<i>attribute 0</i>	<i>attribute 1</i>	<i>attribute 2</i>	...	<i>attribute n</i>
<i>record 0</i>	$w_l^{0,0}$	$w_l^{1,0}$	$w_l^{2,0}$		$w_l^{n,0}$
<i>record 1</i>	$w_l^{0,1}$	$w_l^{1,1}$	$w_l^{2,1}$		$w_l^{n,1}$
...
<i>record m</i>	$w_l^{0,m}$	$w_l^{1,m}$	$w_l^{2,m}$...	$w_l^{n,m}$

Figure 1: Plaintext relation in SQL database

In [Curtmola et. al] the authors formally defined a static index-based symmetric searchable encryption scheme by the following algorithms (see Definition 4.1):

$K \leftarrow \text{Gen}(1^k)$: is a probabilistic key generation algorithm that is run by the user to setup the scheme. It takes as input a security parameter k , and outputs a secret key K .

$(I, \mathbf{C}) \leftarrow \text{Enc}(K, \mathbf{D})$: is a probabilistic algorithm run by the user to encrypt the document collection. It takes as input a secret key K and a document collection $\mathbf{D} = \{D_0, \dots, D_m\}$, and outputs a secure index I and a sequence of ciphertexts $\mathbf{C} = \{C_0, \dots, C_m\}$.

$t \leftarrow \text{Trpdr}(K, w)$: is a deterministic algorithm run by the user to generate a trapdoor for a given keyword. It takes as input a secret key K and a keyword w , and outputs a trapdoor t .

$X \leftarrow \text{Search}(I, t)$: is a deterministic algorithm run by the server to search for the documents in \mathbf{D} that contain a keyword w . It takes as input an encrypted index I for a data collection \mathbf{D} and a trapdoor t and outputs a set X of (lexicographically-ordered) document identifiers.

$D_j \leftarrow \text{Dec}(K, C_j)$: is a deterministic algorithm run by the client to recover a document. It takes as input a secret key K and a ciphertext C_j , and outputs a document D_j .

In order to formally outline Ciphersweet's scheme, we use two cryptographic primitives: a CPA-secure symmetric encryption scheme and a pseudo-random function (PRF). We also use one utility function that performs PRF output truncation. Let's denote $\text{SKE} = (\text{Gen}, \text{Enc}, \text{Dec})$ – CPA-secure symmetric encryption scheme and $f = \{0,1\}^k \times \{0,1\}^x \rightarrow \{0,1\}^y$ – pseudo-random function. Let also $\text{Truncate}(p, \text{value}) \rightarrow \text{value} \upharpoonright_0^p$ be a function that truncates bit vector value to its first p bits. We call “blind index” a truncated PRF output.

Now we are ready to formally define construction of our scheme in context of Curtmola's notion. It is described in Figure 2.

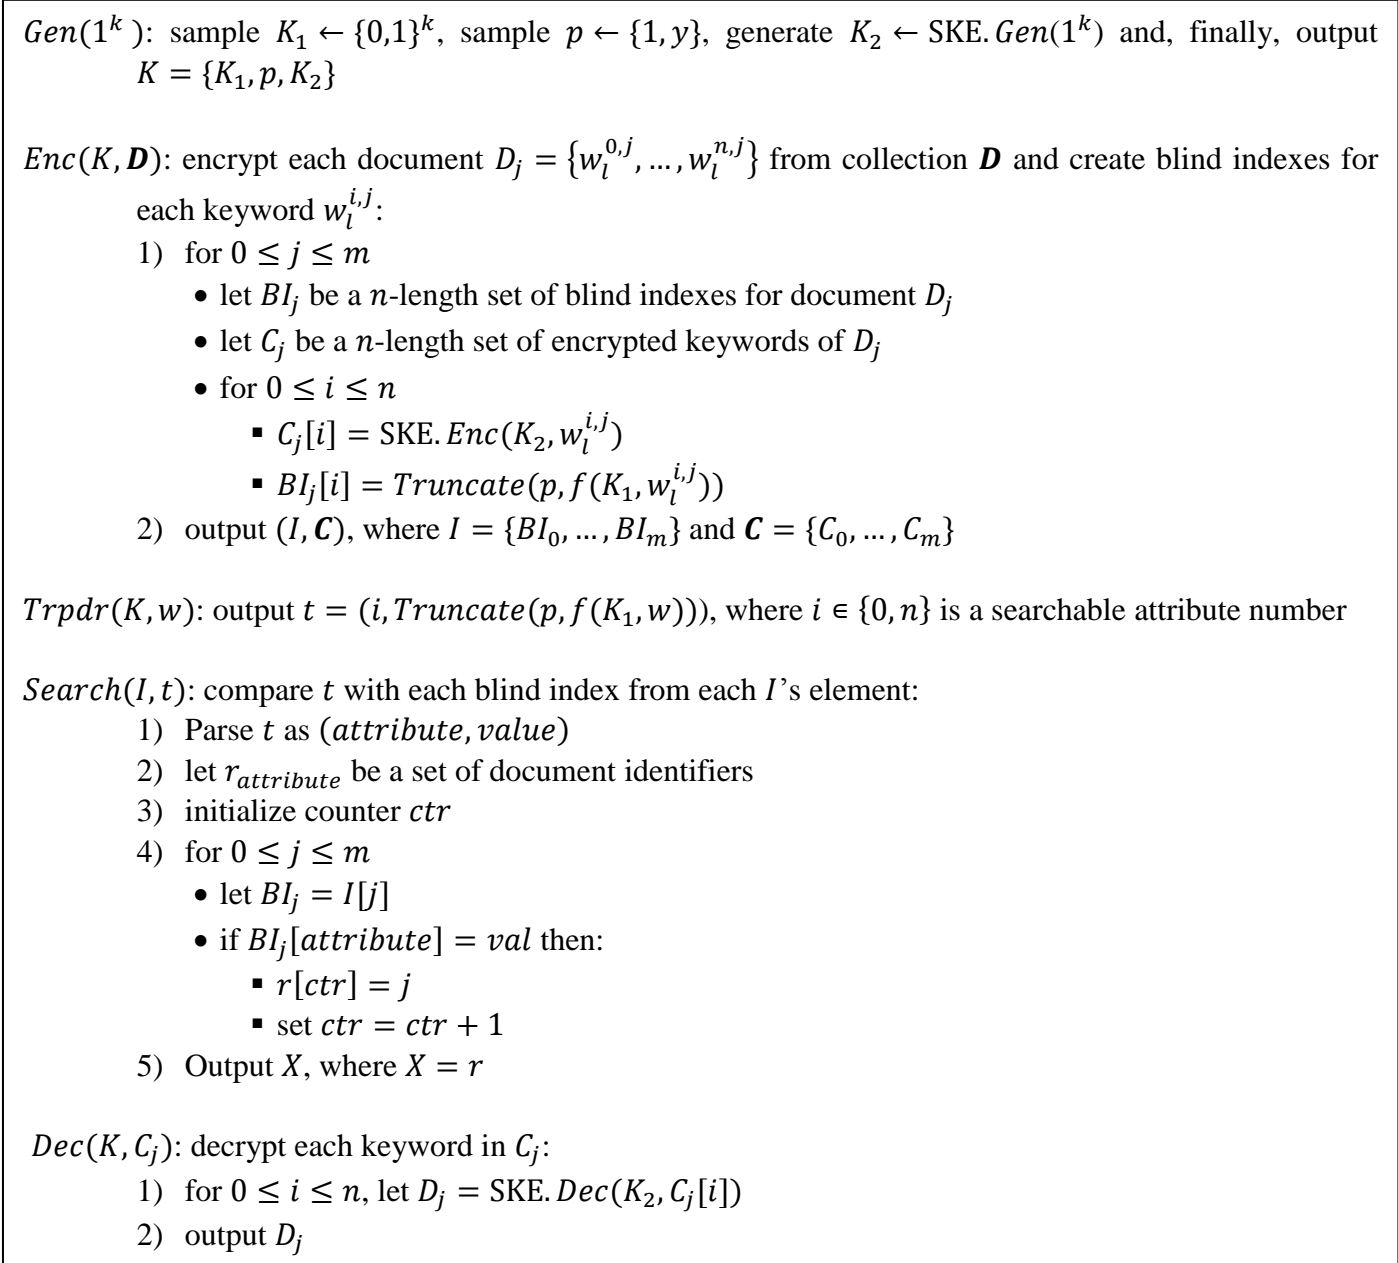


Figure 2: A formal describing of Ciphersweet's scheme

Note that we have to introduce a separate index relation along in order to demonstrate the way, how server should store encrypted collection of documents $\mathbf{C} = \{C_0, \dots, C_m\}$, $C_j = \{c_l^{0,j}, \dots, c_l^{n,j}\}$, and $c_l^{i,j}$ is a single encrypted keyword. Index relation stores $I = \{BI_0, \dots, BI_m\}$ that itself consists from sets of blind indexes $BI_j = \{bi_l^{0,j}, \dots, bi_l^{n,j}\}$ for each encrypted keyword. Figure 3 shows encrypted relation itself, while Figure 4 shows index relation.

	<i>attribute 0</i>	<i>attribute 1</i>	<i>attribute 2</i>	...	<i>attribute n</i>
<i>record 0</i>	$c_l^{0,0}$	$c_l^{1,0}$	$c_l^{2,0}$...	$c_l^{n,0}$
<i>record 1</i>	$c_l^{0,1}$	$c_l^{1,1}$	$c_l^{2,1}$...	$c_l^{n,1}$
...
<i>record m</i>	$c_l^{0,m}$	$c_l^{1,m}$	$c_l^{2,m}$...	$c_l^{n,m}$

Figure 3: Encrypted relation

	<i>attribute 0</i>	<i>attribute 1</i>	<i>attribute 2</i>	...	<i>attribute n</i>
<i>record 0</i>	$bi_l^{0,0}$	$bi_l^{1,0}$	$bi_l^{2,0}$		$bi_l^{n,0}$
<i>record 1</i>	$bi_l^{0,1}$	$bi_l^{1,1}$	$bi_l^{2,1}$		$bi_l^{n,1}$
...
<i>record m</i>	$bi_l^{0,m}$	$bi_l^{1,m}$	$bi_l^{2,m}$...	$bi_l^{n,m}$

Figure 4: Index relation